

Task 1: Data Loading and Initial Inspection

Task Description

Load the data from the V002.json file into a Pandas DataFrame. Create a separate Series for the video URL. Then, display the first 5 rows, the data types of all columns, and the total number of sentences.

Key Pandas Concepts

pd.read_json(), DataFrame/Series creation, .head(), .info(), .shape

```
import pandas as pd
import numpy as np
import json
from collections import Counter

v002_path = "V002.json"

# Load the JSON data
with open(v002_path, 'r', encoding='utf-8') as file:
    raw_data = json.load(file)

# The data is nested under "V002.mp4" key
video_key = "V002.mp4"
video_data = raw_data[video_key]

# Create main DataFrame from instances
df = pd.DataFrame(video_data['instances'])

#todo:

#1. Display first 5 rows
display(df.head())

#2. Display total number of sentences
print(f"Total number of sentences: {df.shape[0]}")

#3. Display columns
print(f"Columns: {list(df.columns)}")
```

		sentence	start	end
split				
0	□□□□□□□□ □□□□□□□□		260	281
train				
1	□□□□□□ □ □□□□□□ □□□□□□□□□□□□ □□ □ □□□□□□ □□□□□...		282	518
train				

```

2  [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] []... 592 742
train
3  [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] []... 1361 1643
train
4  [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] [] []... 1643 1935
test

Total number of sentences: 42
Columns: ['sentence', 'start', 'end', 'split']

```

Task 2: Feature Extraction and Manipulation (Time Features)

Task Description

Convert frame numbers to seconds using video FPS metadata. Create new columns for start time, end time, and duration in seconds.

Key Pandas Concepts

Arithmetic operations on columns, creating new columns, vectorized operations

```

# Task 2 – Convert frame numbers to seconds

# 1. Extract fps from video metadata
fps = video_data["fps"]

# 2. Convert start and end frames to seconds
df["start_sec"] = df["start"] / fps
df["end_sec"] = df["end"] / fps

#todo:
# 3. Calculate duration in seconds (Subtract the end_sec from the
start_sec)
df["duration_sec"] = df["end_sec"] - df["start_sec"]

# 4. Display results
df[["start", "end", "start_sec", "end_sec", "duration_sec"]].head()

```

	start	end	start_sec	end_sec	duration_sec
0	260	281	8.666667	9.366667	0.700000
1	282	518	9.400000	17.266667	7.866667
2	592	742	19.733333	24.733333	5.000000
3	1361	1643	45.366667	54.766667	9.400000
4	1643	1935	54.766667	64.500000	9.733333

Task 3: Data Filtering (Conditional Selection)

Task Description

Create two new DataFrames:

- `df_long`: Contains only sentences that are longer than 5 seconds
- `df_train`: Contains only sentences where the split column is 'train' Report the number of sentences in `df_long` and the average duration of sentences in `df_train`.

Key Pandas Concepts

Boolean masking/conditional filtering, `.loc[]`, `.mean()`, `.count()`

```
# Task 3 – Data Filtering

# 1. Sentences longer than 5 seconds
df_long = df.loc[df["duration_sec"] > 5] #todo

# 2. Sentences belonging to the 'train' split
df_train = df.loc[df["split"] == "train"]

# 3. Reporting
num_long_sentences = df_long.shape[0]
avg_duration_train = df_train["duration_sec"].mean()

print("Number of sentences longer than 5 seconds:",
      num_long_sentences)
print("Average duration of 'train' sentences:", avg_duration_train)

Number of sentences longer than 5 seconds: 31
Average duration of 'train' sentences: 7.031034482758617
```

Task 4: Text Data Preprocessing (String Operations)

Task Description

Clean the sentence column by performing the following steps:

- Remove leading/trailing whitespace
- Check for and count how many sentences contain specific words
- Create a new column `sentence_length` containing the character count of each sentence

Key Pandas Concepts

.str.strip(), .str.contains(), .sum(), .str.len()

Task 4 – Text Data Preprocessing

1. Remove leading/trailing whitespace

```
df['sentence'] = df['sentence'].str.strip()
```

2. Count sentences containing the word "train"

```
count = df['sentence'].str.contains("train").sum()
```

```
print("Number of sentences containing 'train':", count)
```

3. Create a new column for sentence length (character count)

```
df['sentence_length'] = df['sentence'].str.len() #todo
```

4. Preview the updated dataframe

```
df.head()
```

Number of sentences containing 'train': 3

	sentence	start	end
split \			
0	train	260	281
1	train	282	518
2	train	592	742
3	train	1361	1643
4	test	1643	1935

	start_sec	end_sec	duration_sec	sentence_length
0	8.666667	9.366667	0.700000	16
1	9.400000	17.266667	7.866667	133
2	19.733333	24.733333	5.000000	101
3	45.366667	54.766667	9.400000	137
4	54.766667	64.500000	9.733333	117

Task 5: Aggregation and Grouping (Split Analysis)

Task Description

Group the DataFrame by the split column ('train', 'test', 'val'). For each group, calculate and display:

- The total duration of all sentences
- The number of sentences
- The average duration per sentence

Key Pandas Concepts

.groupby(), .agg(), named aggregation

```
# Task 5 – Grouping and Aggregation by 'split'

# Group by the 'split' column and compute aggregations
split_stats = df.groupby('split').agg(
    total_duration=('duration_sec', 'sum'),
    sentence_count=('sentence', 'count'),
    avg_duration=('duration_sec', 'mean')
) #todo

# Display the aggregated results
split_stats
```

	total_duration	sentence_count	avg_duration
split			
test	53.666667	8	6.708333
train	203.900000	29	7.031034
val	38.333333	5	7.666667

Task 6: Missing Value Analysis and Handling

Task Description

In real-world datasets, it is common to encounter missing values. In this task, you will:

- Simulate a missing value in the sentence column by setting the first row to None
- Check how many missing values exist in the sentence column
- Replace missing values with the placeholder string: "[MISSING SENTENCE]"

Key Pandas Concepts

.loc[] for value assignment, .isnull(), .sum(), .fillna()

Task 6 – Missing Value Analysis and Handling

1. Simulate a missing value in the first row
df.loc[0, "sentence"] = None

2. Count missing values in the 'sentence' column
missing_count = df['sentence'].isnull().sum() #todo
print("Number of missing sentences:", missing_count)

3. Replace missing values with '[MISSING SENTENCE]'
df['sentence'] = df['sentence'].fillna("[MISSING SENTENCE]")

4. Preview the updated DataFrame
df.head()

Number of missing sentences: 0

	sentence	start	end
split \			
0	□□□□□□□□ □□□□□□□□	260	281
train			
1	□□□□□□ □ □□□□□□ □□□□□□□□□□□□ □□ □ □□□□□□ □□□□□...	282	518
train			
2	□□ □□□□□□ □□□□□□□ □□□□□□ □□□□□□□□□□□□ □□□□□□□□...	592	742
train			
3	□□□□ □□□□□ □□□□□□□□□ □□□□□□□□□ □□□□□□ □□□□□□□□...	1361	1643
train			
4	□□ □□□□ □□□□□□ □□□□□□ □□□□□□ □□□□□□□□ □□ □□□□□...	1643	1935
test			

	start_sec	end_sec	duration_sec	sentence_length
0	8.666667	9.366667	0.700000	16
1	9.400000	17.266667	7.866667	133
2	19.733333	24.733333	5.000000	101
3	45.366667	54.766667	9.400000	137
4	54.766667	64.500000	9.733333	117

Task 7: Indexing and Hierarchical Data Access

Task Description

In this task, you will practice indexing and hierarchical data access in Pandas.

- Set the split column as the DataFrame's index

- Use the .loc[] accessor to select specific rows/columns:
 - All rows belonging to the 'test' split
 - Only start_sec and end_sec columns for rows in the 'val' split

Key Pandas Concepts

.set_index(), Hierarchical indexing, .loc[] for row and column selection

Task 7 – Indexing and Hierarchical Data Access

1. Set 'split' column as the index

```
df_indexed = df.set_index("split")
```

2. Select all rows belonging to the 'test' split

```
df_test = df_indexed.loc["test"]
```

```
print("All rows in 'test' split:")
```

```
display(df_test)
```

3. Select only start_sec and end_sec for the 'val' split

```
df_val_times = df_indexed.loc["val", ["start_sec", "end_sec"]] #todo
```

```
print("\nStart and End times for 'val' split:")
```

```
display(df_val_times)
```

All rows in 'test' split:

		sentence	start	end
\	split			
test		00 0000 000000 000000 000000 00000000 00 00000...	1643	1935
test		000000000000 000000000 00 00000000 000000 0000...	6995	
7228				
test		00000 00000 000000 00000000000000 00000000 0000	7777	7886
test		0000 0000000000 00000000 000000000000 0 000000 ...	12404	
12597				
test		00000 0000000 0000 0000 0000 00000000000 00 000...	14993	15212
test		00000 0000000 0000000000 0000000 0000000 00000...	15708	
15928				
test		00000 00000 00000 0000 0000 000 000 000 19183	19314	
test		00000000 000000 00000000 000000 000 00 0000 00...	20944	21157
	start_sec	end_sec	duration_sec	sentence_length
split				
test	54.766667	64.500000	9.733333	117
test	233.166667	240.933333	7.766667	121

test	259.233333	262.866667	3.633333	46
test	413.466667	419.900000	6.433333	86
test	499.766667	507.066667	7.300000	97
test	523.600000	530.933333	7.333333	81
test	639.433333	643.800000	4.366667	39
test	698.133333	705.233333	7.100000	71

Start and End times for 'val' split:

	start_sec	end_sec
split		
val	240.933333	247.466667
val	247.466667	259.233333
val	392.500000	396.066667
val	476.233333	487.000000
val	656.800000	662.500000

Task 8: Handling Categorical Data

Task Description

In this task, you will optimize the DataFrame for memory efficiency by converting the split column into a categorical data type.

Steps:

- Check the memory usage of the DataFrame before conversion
- Convert the split column to Categorical using `.astype('category')`
- Check the memory usage after conversion
- Understand why converting string columns to categorical types is useful in ML

Key Pandas Concepts

`.astype('category')`, `.memory_usage()`, Data type optimization

```
# Task 8 – Handling Categorical Data

# 1. Check memory usage before conversion
memory_before = df.memory_usage(deep=True).sum()
print("Memory usage before conversion:", memory_before, "bytes")

# 2. Convert 'split' column to categorical
df['split'] = df['split'].astype('category')

# 3. Check memory usage after conversion
memory_after = df.memory_usage(deep=True).sum()
```



```
print("Memory usage after conversion:", memory_after, "bytes")

# 4. Display the DataFrame info to verify
df.info()

Memory usage before conversion: 15414 bytes
Memory usage after conversion: 13473 bytes
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42 entries, 0 to 41
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sentence              42 non-null    object
1   start                 42 non-null    int64
2   end                  42 non-null    int64
3   split                42 non-null    category
4   start_sec            42 non-null    float64
5   end_sec              42 non-null    float64
6   duration_sec         42 non-null    float64
7   sentence_length      42 non-null    int64
dtypes: category(1), float64(3), int64(3), object(1)
memory usage: 2.6+ KB
```

Task 9: Combining Data (Merging and Concatenation)

Task Description

In this task, you will combine sentence-level data with video-level metadata.

Steps:

- Create a separate DataFrame `df_info` containing video-level metadata: url, date, presenter, and signer
- Add relevant metadata columns (date and presenter) to the sentence DataFrame using broadcasting

Key Pandas Concepts

DataFrame creation from scalar values, Broadcasting, `pd.merge()` (conceptually)

```
# Task 9 – Combining Data

# 1. Create a separate DataFrame for video-level metadata
df_info = pd.DataFrame({
    "url": [video_data["url"]],
```

```

    "date": [video_data["date"]],
    "presenter": [video_data["presenter"]],
    "signer": [video_data["signer"]]
})

```

2. Add date and presenter to the main sentence DataFrame using broadcasting

```

df['date'] = video_data["date"]
df['presenter'] = video_data["presenter"]
df['signer'] = video_data["signer"]
#todo

```

3. Preview the updated DataFrame

```
df.head()
```

		sentence	start	end
split \				
0		□□□□□□□□ □□□□□□□□	260	281
train				
1	□□□□□□ □ □□□□□□ □□□□□□□□□□□□ □□ □ □□□□□□ □□□□□...		282	518
train				
2	□□ □□□□□□ □□□□□□□ □□□□□□ □□□□□□□□□□□□ □□□□□□□□...		592	742
train				
3	□□□□ □□□□□ □□□□□□□□□□ □□□□□□□□□ □□□□□ □□□□□□□□...		1361	1643
train				
4	□□ □□□□ □□□□□□ □□□□□□ □□□□□□ □□□□□□□□ □□ □□□□□...		1643	1935
test				

	start_sec	end_sec	duration_sec	sentence_length		date	\
0	8.666667	9.366667	0.700000	16	8	DEC	2017
1	9.400000	17.266667	7.866667	133	8	DEC	2017
2	19.733333	24.733333	5.000000	101	8	DEC	2017
3	45.366667	54.766667	9.400000	137	8	DEC	2017
4	54.766667	64.500000	9.733333	117	8	DEC	2017

	presenter	signer
0	□□□□□□□□ □□□□□□	Shahnaz Rita
1	□□□□□□□□ □□□□□□	Shahnaz Rita
2	□□□□□□□□ □□□□□□	Shahnaz Rita
3	□□□□□□□□ □□□□□□	Shahnaz Rita
4	□□□□□□□□ □□□□□□	Shahnaz Rita

Task 10: Data Visualization Preparation (Value Counts & Binning)

Task Description

In this task, you will prepare the dataset for analysis and visualization by performing the following steps:

- Determine the top 3 most frequent words in the sentence column
- Create a new column `duration_bin` by grouping the `duration_sec` into three bins: Short (0-5s), Medium (5-10s), and Long (>10s)
- Use `.value_counts()` on the new `duration_bin` column to show the distribution of sentence lengths

Key Pandas Concepts

`.str.lower()`, `.str.split()`, value counting, `pd.cut()` for binning

Task 10 – Data Visualization Preparation

```
from collections import Counter
```

1. Determine top 3 most frequent words

```
all_words = df['sentence'].str.lower().str.split().sum() # Flatten list of all words
```

```
word_counts = Counter(all_words)
```

```
top_3_words = word_counts.most_common(3)
```

```
print("Top 3 most frequent words:", top_3_words)
```

2. Create duration bins

```
bins = [0, 5, 10, df['duration_sec'].max()]
```

```
labels = ['Short', 'Medium', 'Long']
```

```
df['duration_bin'] = pd.cut(df['duration_sec'], bins=bins, labels=labels, right=False)
```

3. Show distribution of sentence lengths by duration_bin

```
duration_distribution = df['duration_bin'].value_counts()
```

```
print("\nDistribution of sentence lengths:")
```

```
print(duration_distribution)
```

```
Top 3 most frequent words: [(' ', 17), ('the', 9), ('a', 8)]
```

```
Distribution of sentence lengths:
```

```
duration_bin
```

```
Medium      25
```

```
Short       10
```

```
Long      6  
Name: count, dtype: int64
```

Summary

Tasks Completed:

1. Data Loading and Initial Inspection
2. Feature Extraction and Manipulation (Time Features)
3. Data Filtering (Conditional Selection)
4. Text Data Preprocessing (String Operations)
5. Aggregation and Grouping (Split Analysis)
6. Missing Value Analysis and Handling
7. Indexing and Hierarchical Data Access
8. Handling Categorical Data
9. Combining Data (Merging and Concatenation)
10. Data Visualization Preparation (Value Counts & Binning)

All tasks have been successfully completed using core Pandas operations essential for Machine Learning preprocessing.